

CS 273P: Machine Learning and Data Mining

Homework 2

Due date: See Canvas

Instructor: Xiaohui Xie

This homework is graded automatically on **Gradescope** using an autograder. All questions are designed to be **unit-testable**: you will implement specific functions that return the specified results. **No plots and no PDF writeup are required.**

Allowed libraries: You should use the standard Python scientific stack: **numpy**, **scipy**, **pandas**, and **scikit-learn**.

Submission rules (READ CAREFULLY):

1. Download the starter files from Canvas.
2. Complete your work by editing only the provided Python files: **problem1.py** and **problem2.py**.
3. Create a plain text file called **statement.txt** containing your collaboration statement.
4. Submit to Gradescope by uploading a single **ZIP file** containing exactly:
 - **problem1.py**
 - **problem2.py**
 - **statement.txt**

No folders/subdirectories; all files must be at the top level of the zip.

5. You may resubmit before the deadline; your **highest score** counts.

If you run into issues with the autograder or submission format, post on Ed Discussion so others can benefit.

Summary: (1) Edit **problem1.py**, **problem2.py**; (2) Add **statement.txt**; (3) Zip all three files and upload to Gradescope; (4) Autograder score determines your grade.

Points: This homework adds up to a total of **100 points**, as follows:

Problem 1: Linear/Polynomial Regression + Model Selection	60 points
Problem 2: Cross-Validation	35 points
Statement of Collaboration	5 points

Problem 1: Linear/Polynomial Regression + Model Selection (60 points)

In this problem we explore regression, feature construction, rescaling, and model selection. You may use **numpy** for array operations and **scikit-learn** utilities/models where appropriate.

Dataset. Load the dataset **data/curve80.txt**. The first column is x and the second column is y . **Do not shuffle the data.** Split into 75% train and 25% test by taking the first 75% as training and the remaining as test.

All required functions for Problem 1 are in **problem1.py**. The autograder calls them directly.

1. **Data loading and split (10 points).** Implement:

- **load_curve80(path) -> (X, y):** return $X \in R^{N \times 1}$ and $y \in R^N$.
- **split_data(X, y, frac=0.75) -> (Xtr, Xte, ytr, yte):** deterministic split without shuffling.
- **shapes(Xtr, Xte, ytr, yte) -> (sxtr, sxte, sytr, syte):** return the shapes as tuples.

2. **Baseline linear regression (20 points).** Fit a degree-1 regression model (with intercept) on the training set and evaluate MSE. Implement:
- `fit_linear(Xtr, ytr) -> model`: fit a linear regression model with intercept (you may use `sklearn.linear_model.LinearRegression`).
 - `predict(model, X) -> yhat`: return predictions as a 1D array.
 - `mse(yhat, y) -> float`: return mean squared error.
 - `eval_linear(Xtr, ytr, Xte, yte) -> (mse_tr, mse_te)`.
3. **Polynomial regression with scaling (20 points).** Train polynomial regression models by expanding x to polynomial features and standardizing the features. You may use `sklearn.preprocessing.PolynomialFeatures` and `sklearn.preprocessing.StandardScaler`. Implement:
- `make_poly_pipeline(degree = 3) -> model`: return a fitted or configurable model/pipeline that (i) creates polynomial features up to `degree`, set `degree` to 3 for this problem, (ii) standardizes features, and (iii) performs linear regression with intercept. (*Hint: Pipeline(PolynomialFeatures, StandardScaler, LinearRegression)*.)
 - `fit_poly(Xtr, ytr, degree = 3) -> model`: fit the pipeline on training data.
 - `eval_poly(Xtr, ytr, Xte, yte, degree = 3) -> (mse_tr, mse_te)`.
4. **Degree sweep + recommendation (10 points).** Train polynomial models for degrees $d \in \{1, 3, 5, 7, 10, 18\}$ and compute train/test MSE for each. Implement:
- `eval_degrees(Xtr, ytr, Xte, yte, degrees) -> (mse_tr, mse_te)` returning two 1D arrays aligned with `degrees`.
 - `recommend_degree(degrees, mse_te) -> int` returning the degree with the smallest test MSE; break ties by choosing the smaller degree.

Problem 2: Cross-Validation (35 points)

In Problem 1 you selected a polynomial degree using a held-out test set. Now assume you do **not** use the test labels for selection, and instead use K -fold cross-validation on the training set (X_{tr}, y_{tr}) .

All required functions for Problem 2 are in `problem2.py`.

1. **K -fold splitting (10 points).** Implement:
- `kfold_indices(n, K) -> list_of_folds`: deterministically split indices $\{0, \dots, n - 1\}$ into K folds in order (no shuffling). Return a list of length K , each an integer array of validation indices.
 - `train_val_split(X, y, folds, i) -> (Xti, yti, Xvi, yvi)`: use fold i as validation and the rest as training.
2. **K -fold CV error (15 points).** Compute the average validation MSE of polynomial regression for a given degree. You may reuse your pipeline from Problem 1. Implement:
- `cv_mse_poly(Xtr, ytr, degree = 3, K) -> float`: perform K -fold CV on (X_{tr}, y_{tr}) ; on each fold fit the polynomial+scaling+linear-regression pipeline using only the fold's training split, then compute validation MSE; return the average across folds.
3. **Degree selection by CV (10 points).** Using degrees $d \in \{1, 3, 5, 7, 10, 18\}$:
- `cv_curve(Xtr, ytr, degrees, K=5) -> cv_mses`: return a 1D array of CV MSEs aligned with `degrees`.
 - `recommend_degree_cv(degrees, cv_mses) -> int`: return the degree with smallest CV MSE; break ties by choosing the smaller degree.

Statement of Collaboration (5 points)

Add a plain text file `statement.txt` to your submission. List the names of collaborators and the nature of collaboration, or write “No collaboration.” if you worked alone. Do not share code. Academic honesty policies apply.